

Homework 3:

Representation Learning & Recommender Systems

Mobile Data Mining

Spring 2019

Goal

- **Representation Learning**

- Use representation learning techniques in mobile big data mining
 - Learning the embeddings of apps based on the app usage traces
- Familiar with some simple classification algorithms
 - Classify apps into different categories using the obtained embeddings as the feature

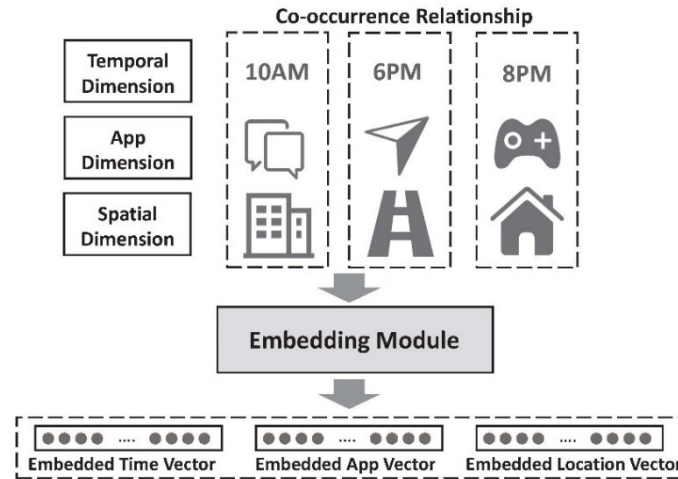
- **Recommender Systems**

- Use collaborative filtering to recommend apps to users
 - User-user collaborative filtering
 - Item-item collaborative filtering
 - Matrix-factorization (MF)-based collaborative filtering

Data

- code/input/AppUsageTrace.txt
 - App usage traces in Shanghai
 - Involve over 2,000 apps and 10,875 locations (cellular base stations)
 - Duration about one week
 - Format (each line)
 - User ID || Location ID || Time (in hour) || ID of Used App
- code/App2Category.txt
 - Category of each app
 - Format (each line)
 - App ID || Category ID
- code/Categories.txt
 - English name of each category
 - Format (each line)
 - Category ID || English Name

Representation Learning: Reconstruction embedding



- Utilize the co-occurrence relationship between units (include time-bins, locations, and apps) to learning their embeddings
- Optimization Target:

$$\mathcal{O} = - \sum_{r \in R} \sum_{e \in r} \log P(e | r_{-e})$$

损失函数

- R is the set of all app usage records. Each record contains 3 units, i.e., the time-bin, the base station, and the used app. In addition, $r_{-e} = \{o \neq e | o \in r\}$.

Representation Learning: Reconstruction embedding

- We model the likelihood of observing each unit e in each record r given its context r_{-e} by

$$P(e|r_{-e}) = \exp(s(e, r_{-e})) / \sum_{o \in X} \exp(s(o, r_{-e}))$$

- X represents all with the same type of e . For example, for an app e , X is the set of all apps.
- $s()$ is a score function reflecting the similarity between the unit e and its context r_{-e} , defined by:

$$s(e, r_{-e}) = \langle I_e, h_e \rangle$$
$$h_e = \frac{1}{|r_{-e}|} \sum_{o \in r_{-e}} I_o$$

- where $I_e \in \mathcal{R}^D$ is the D-dimensional embedding of unit e .

Representation Learning:

Reconstruction embedding

Negative Sampling:

- Calculating $P(e|r_{-e})$ requires the summation over the entire set of units X , which leads to high computational complexity

- We approximate $-\log P(e|r_{-e})$ with

$$J_r = -\log \sigma(s(e, r_{-e})) - \sum_{k=1}^K \log \sigma(-s(o_k, r_{-e}))$$

- where $o_1 \dots o_K$ are the K negative samples randomly sampled from X
- $\sigma(x) = \frac{\exp(x)}{1+\exp(x)}$ is the sigmoid function.

Representation Learning: Reconstruction embedding

Stochastic gradient descent:

- Iteratively randomly sample a record r and a unit e . Then change variables in the direction of the gradient of J_r

- For example, we have

$$\frac{\partial J_r}{\partial I_e} = (\sigma(s(e, r_{-e})) - 1)h_e$$

- Thus, each time we will update I_e with

$$I_e \leftarrow I_e - \alpha \frac{\partial J_r}{\partial I_e}$$

- where α is the learning rate.

Recommender Systems: Preprocessing

- Construct the user-app matrix $R = \{r_{ul}\}_{U \times L}$
 - U is the set of all users, and L is the set of all apps.
 - r_{ul} represents whether user u has used app l .

Recommender Systems: User-User Collaborative Filtering

- Predict r_{ul} by

$$r_{ul} = \frac{\sum_{v \in N(u)} S_{uv} r_{vl}}{\sum_{y \in N(u)} S_{uy}}$$

- where $S_{uv} = \text{sim}(u, v)$, $N(u)$ is the set of k users most similar to u .

- Cosine similarity measure

$$\text{sim}(u, v) = \cos(\mathbf{r}_u, \mathbf{r}_v) = \frac{\mathbf{r}_u \cdot \mathbf{r}_v}{\|\mathbf{r}_u\| \cdot \|\mathbf{r}_v\|}$$

- where \mathbf{r}_u is the u th row in R .

Recommender Systems: Item-Item Collaborative Filtering

- Predict r_{ul} by

$$r_{ul} = \frac{\sum_{j \in N(l)} S_{lj} r_{uj}}{\sum_{j \in N(l)} S_{lj}}$$

- where $S_{lj} = \text{sim}(l, j)$, $N(l)$ is the set of k apps most similar to l .

- Cosine similarity measure

$$\text{sim}(l, j) = \cos(\mathbf{r}_l, \mathbf{r}_j)$$

- where \mathbf{r}_l is the l th column in R .

Recommender Systems: Matrix-Factorization (MF)-based Collaborative Filtering

- Predict r_{ul} by

$$r_{ul} = q_u \cdot p_l$$

- where q_u and p_l are k -dimensional latent vectors for user u and app l , respectively

- q_u and p_l are learnt from solving the optimization problem

$$\min_{\{P, Q\}} \sum_{\text{Training set}} (r_{ul} - q_u \cdot p_l)^2 + \lambda_1 \sum_u ||p_u||^2 + \lambda_2 \sum_l ||q_l||^2$$

- which can be solved by stochastic gradient descent (SGD).
- For detailed algorithm, refer to page 79 of the slides of recommender systems in course documents.

Experiments - 1

- Embedding

- A line in “embed.py” is missing. Please complete it.
- Use the code to learn the embedding of apps
- Run: “python train.py”
- Output: code/output/embeddings/

- Virtualization

- Use the code “DimReduce.py” to map all the obtained embedding into 2-dimensional space for virtualization
- Plot the embeddings of different time-bins (hours)
- Plot the embeddings of apps of the category “Video”, “Finance”, “Music”
- Bonus: Tune the parameter ‘negative’ and ‘dim’, try to find some interesting observations.

- Classification

- Focus on apps of “video” and “Finance”
- Use SVM or Decision Trees to divide them based on their embeddings
 - You can directly call existing functions of these algorithms
 - Use 70% as the training set and 30% apps the test set
- Calculate the obtained accuracy

Experiments - 2

- Randomly remove 20% user-app pair in the dataset, and use collaborative filtering to predict them
 - prediction whether corresponding users have used corresponding apps
 - Calculate the RMSE
 - $RMSE = \sqrt{\frac{1}{n} \sum (r_{ui} - \hat{r}_{ui})^2}$
 - \hat{r}_{ui} is the predicted value, n is the number of predicted user-app pair.
 - Implement user-user, item-item, MF-based collaborative filtering.
 - Investigate the influence of k . Plot the RMSE- k curve.

Submission

- Submit this homework before June 16th. (Hard Deadline, please keep in mind)
- Submit as **.zip** file, including:
 - 1) A word document, Including:
 - Brief summary about the algorithms
 - All the results you obtained, presented in table or figure (using figure more, and show the results clearly and beautifully)
 - Interpretation/discussion for each result
 - Do not need to copy the code into this document
 - 2) Source code, Including:
 - Data processing code
 - Collaborative filtering code
 - Other analysis code

Thank you!